


Performance of Distributed and Shared Memory Parallel Sparse Matrix Vector Multiplication

Kristian Sordal



Thesis for Master of Science Degree at the University
of Bergen, Norway

2025

©Copyright Kristian Sordal

The material in this publication is protected by copyright law.

Year: 2025

Title: Performance of Distributed and Shared Memory
Parallel Sparse Matrix Vector Multiplication

Author: Kristian Sordal

Acknowledgements

Abstract

SPARSE MATRIX VECTOR MULTIPLICATION is an important kernel used in scientific computing. It is a problem that lends itself well to parallelization. The problem is bounded by, and scales with the memory bandwidth of the system. Therefore in order to efficiently perform *SpMV* on large distributed memory systems, it is important to reduce the communication between nodes, in order to extract as much as possible out of the memory bandwidth of the system.

This thesis aims to investigate the results of *SpMV* when ran using different communication strategies.

Contents

Acknowledgements	iii
Abstract	v
1 Introduction	1
1.1 Research Question	1
2 Theory	3
2.1 Sparse Matrix-Vector Multiplication	3
2.2 Definitions	3
2.3 Amdahl's Law	4
2.4 Latency	4
3 Background	7
3.1 CSR Storage Format	7
3.1.1 Computational Intensity	8
3.2 Other storage formats	8
3.2.1 COO Format	8
3.2.2 CSC Format	9
3.2.3 ELLPack Format	9
3.3 Sequential SpMV	10
3.4 Shared Memory SpMV	11
3.4.1 Scheduling options	12
3.4.2 Dynamic scheduling	12
3.4.3 First touch policy	13
3.4.4 Performance impact of first touch policy	13
3.5 Distributed Memory SpMV	13
3.5.1 Load Balancing	14
4 Communication Strategies	15
4.1 Exchange entire vector	16

4.2	Exchange only separators	16
4.2.1	Reordering	17
4.3	Exchange only required separators	18
4.4	Exchange only required separator values	19
4.5	2d - Fully scalable communication	19
5	Results	21
6	Previous Work	25

Chapter 1

Introduction

1.1 Research Question

This thesis aims to investigate the impact different communication strategies have on the performance of parallel SpMV. The implementation of SpMV uses a shared memory layout on each socket of a node, and distributed memory for communication across nodes.

Chapter 2

Theory

2.1 Sparse Matrix-Vector Multiplication

Sparse Matrix-Vector Multiplication (SpMV) is a fundamental operation encountered in many areas of scientific computing. It is especially prominent in solving large systems of linear equations and in large-scale simulations. The matrices involved are typically both very large and very sparse.

A matrix can in theory be considered sparse if it is worthwhile to treat zero values separately. In theory, this translates to a matrix being less than full, i.e. less than $\mathcal{O}(n^2)$ nonzeros for a $n \times n$ matrix. However, in the context of sparse linear algebra, sparse means that there is a constant number of nonzeros per row, i.e. $\mathcal{O}(n)$ nonzeros per row. The matrices used in scientific computing, such as matrices based on meshes, or graphs such as social networks all have this property. Optimizing the performance of SpMV, particularly through parallel computing techniques, is crucial for enhancing the efficiency of many scientific applications.

However, SpMV is notoriously difficult to optimize, both in sequential and parallel implementations. One major reason is its inherently low computational intensity.

2.2 Definitions

Definition 2.2.1 (Separator). In the context of SpMV, a separator is a node in the graph that has an edge that strides between two partitions.

2.3 Amdahl's Law

Amdahl's Law provides a theoretical framework for understanding the limits of performance improvement when additional computational resources are applied to a given problem. It quantifies the potential speedup achieved by optimizing a specific portion of a system, emphasizing that the overall gain is constrained by the proportion of time the optimized component contributes to execution.

Definition 2.3.1 (Amdahl's Law). The maximum achievable speedup of a computation is limited by the fraction of execution time that remains sequential, even when an arbitrarily large number of parallel resources is employed.

In the context of parallel computing, this principle highlights that while increasing the number of processing units can accelerate the parallelizable portion of a workload, the sequential fraction imposes a fundamental performance ceiling. Formally, if S denotes the total speedup, t_p represents the fraction of execution time that can be parallelized, and s_p is the speedup achieved for that parallelizable portion, Amdahl's Law is expressed as:

$$S = \frac{1}{(1 - t_p) + \frac{t_p}{s_p}} \quad (2.1)$$

This equation reveals that as $s_p \rightarrow \infty$, the theoretical maximum speedup approaches $\frac{1}{1 - t_p}$, illustrating that the non-parallelizable portion becomes the dominant limiting factor in scalability.

2.4 Latency

It is worthwhile to speak on the typical latency numbers for various operations on a computer.

Operation	Time [ns]
L1 cache reference	1
Branch misprediction	3
L2 cache reference	4
Mutex lock/unlock	17
Main memory reference	100
Compress 1K bytes with Zippy	2000
Send 2kB over 10 Gbps network	1600
Send 1K bytes over 1 Gbps network	10 000
Read 4K bytes randomly from SSD*	20 000
Round trip within same datacenter	500 000
Read 1MB sequentially from memory	1 000 000
Disk seek	10 000 000
Read 1MB sequentially form disk	10 000 000
TCP packet round trip between continents	150 000 000

Figure 2.1: Latency numbers for common operation, adapted from [2]

Chapter 3

Background

3.1 CSR Storage Format

CSR (Compressed Sparse Row) is the most widely used storage format for sparse matrices. As its name suggests, it compresses the amount of memory used to store a matrix without loss of information. It does so by utilizing three vectors A_p, A_j, A_x . Figure 3.1 shows an example of a matrix stored in CSR format, adapted from [1].




Figure 3.1: Example matrix represented in CSR Format.

The first vector, A_p stores the indices of the first non-zero in the vectors A_p and A_x . For a given entry $A_p[i]$, $A_p[i]$ is the index of the first non-zero in the i^{th} row. $A_j[j]$ and $A_x[j]$ denotes the column index and value of the j^{th} non-zero, respectively.

Throughout the remainder of this thesis, we operate under the assumption that all matrices are represented in CSR (Compressed Sparse Row) format, unless explicitly noted otherwise.

3.1.1 Computational Intensity

The *computational intensity* of an operation describes the relation between the number of floating-point operations (FLOPS) and the number of memory accesses required. It is formally defined as:

$$\text{Computational intensity} = \frac{\text{FLOPS}}{\text{Memory accesses}} \quad (3.1)$$

Operations with low computational intensity, such as SpMV, are often *memory bound* rather than *compute bound*. This means that increasing the computational power of a system (e.g., faster processors) does not necessarily lead to proportional speedups in SpMV performance, as memory bandwidth remains the limiting factor.

3.2 Other storage formats

There are of course other ways to store a matrix, each with its own benefits and drawbacks.

3.2.1 COO Format

The Coordinate List (COO) format stores the matrix as a set of triples on the form i, j, x , where i is the row index, j is the column index, and x is the value stored at (i, j) in the matrix. In this format all 0 entries are ignored.




Figure 3.2: cooformat

3.2.2 CSC Format

The Compressed Sparse Column (CSC) format is similar to the CSR format, but instead of compressing the rows, we compress the columns.




Figure 3.3: cscformat

3.2.3 ELLPack Format

For an $M \times N$ matrix with a maximum of K non-zeros per row, the ELLPack format stores the non-zeros in an $M \times K$ matrix `data`, and

an $M \times K$ matrix **indices**. The **data** matrix store the values of the non-zeros, and **indices** store the column index of every element. Rows that have fewer than K non-zeros are padded with zeros. Adapted from [3].




Figure 3.4: ellpackformat

3.3 Sequential SpMV

A sequential implementation of SpMV on a matrix stored in the CSR format can be implemented in the following manner:

Algorithm 1: Sequential CSR-based SpMV

Input : A_p, A_j, A_x, x
Output : y

```

for  $i \leftarrow 0$  to  $n$  do
     $\text{sum} \leftarrow 0$ 
    for  $j \leftarrow A_p[i]$  to  $A_p[i+1]$  do
         $\text{sum} = \text{sum} + A_x[j] \cdot x[A_j[j]]$ 
     $y[i] \leftarrow \text{sum}$ 

```

For SpMV on well structured matrices, i.e. those similar to the matrix shown in Figure 3.5 we read 12 bytes, and perform 2 FLOPS for each

non-zero in the matrix. Keen-eyed readers might notice that this does not coincide with the amount of FLOPS read per non-zero in algorithm 3.3. Here we read two doubles, and one integer, which would be equivalent to 20 bytes. The reason for the discrepancy is due to the fact that after the first time x is accessed, it is loaded into cache, and heavily reused in subsequent iterations, and can for that reason be disregarded.

For heavily unstructured matrices, it is possible that we actually read up to 76 bytes per 2 FLOPS. This will occur if there is no cache reuse, and a new cache line (64 bytes) is loaded for each non-zero.




Figure 3.5: Well structured (a) and poorly structured (b) matrices.

3.4 Shared Memory SpMV

SpMV can be parallelized using the OpenMP directive `#pragma omp parallel for`. By default, this tells OpenMP to use `static` scheduling when parallelizing the outer iteration loop. When static scheduling is used, the span of the iteration that each thread will execute is precomputed, and stays static, as the naming suggests. There are other scheduling options, such as `dynamic` and `guided`, which will be discussed in later sections.

An implementation of shared memory SpMV is outlined below.

Algorithm 2: Shared Memory CSR-based SpMV

Input : A_p, A_j, A_x, x
Output : y

```

#pragma omp parallel for
for  $i \leftarrow 0$  to  $n$  do
  sum  $\leftarrow 0$ 
  for  $j \leftarrow A_p[i]$  to  $A_p[i+1]$  do
    sum = sum +  $A_x[j] \cdot x[A_j[j]]$ 
   $y[i] \leftarrow$  sum
  
```

3.4.1 Scheduling options

As seen in algorithm 2, the outer loop is parallelized, which translates to dividing the rows of the matrix evenly among the threads. This works fine for well structured matrices, but for matrices with dense rows, such as the matrix shown in Figure 3.6, we obtain large imbalances in the computational load for each thread, which impacts performance.




Figure 3.6: staticscheduling

3.4.2 Dynamic scheduling

With dynamic scheduling, OpenMP still precomputes the span of iterations, however threads are not assigned a specific set of iterations to execute. If thread A is assigned a set of sparse rows, and thread B is assigned a set of dense rows, then with static scheduling, A will be idle and

not perform any computation while B is busy with its dense rows. With dynamic scheduling, iterations will be dynamically assigned to threads that are idle, which can distribute the workload more evenly in cases where the computational load of each iteration can differ.

At first glance, it might seem that this will fix the problem static scheduling poses in regards to workload imbalance. This can be the case on single socketed machines with few physical cores such as on a personal laptop, but as the systems that are used in this thesis are dual-socketed, and have a large number of physical cores, it is crucial to consider the impact that *first touch policy* has.

3.4.3 First touch policy

In short, first touch policy means that the first thread to access a memory page will be the one to allocate it to its local memory. This means that if thread A accesses a memory page, and then thread B accesses the same memory page, thread B will not have access to the data in its local memory, and will have to access it from the main memory. This can lead to performance degradation if the data is accessed frequently, as accessing data from main memory is significantly slower than accessing it from local memory.

3.4.4 Performance impact of first touch policy

As seen in table 2.1, accesses to main are significantly slower than accesses to say the L1 cache, and on dual-socketed nodes, accesses to main memory on the opposite socket have to be sent through the interlink, which is once again significantly slower than accessing main memory. From this it becomes evident that using dynamic (or guided) scheduling is not a solution to the performance issues that might arise from poorly structured matrices.

3.5 Distributed Memory SpMV

In distributed memory, the workload is spread across multiple nodes. Physically, each node is its own dual-socketed system, each with its own memory. If data stored on one node is needed by another node, explicit communication has to occur between the nodes. This is typically done using the Message Passing Interface (MPI).

For distributed memory SpMV, algorithm 2 is still utilized for computing the result vector y , but the matrix is split between nodes, and each node computes a local part of y . The final vector has to be assembled at the end of each iteration of SpMV.

On distributed memory systems, the challenges discussed in 3.4.1 are still relevant, but such systems also pose challenges in regards to the communication volume between nodes. As is evident from Table 2.1, communication between nodes (i.e. sending bytes over the network) is slow. As communication between nodes is slow, reducing the total communication volume per iteration of SpMV is crucial for performance.

3.5.1 Load Balancing

In order to reduce the total communication volume, it is necessary to partition the graph into several part, usually one part per node, or one part per socket. The question falls on how to most effectively partition the graph in order to reduce the communication volume. Firstly, the graph partitioner needs to ensure that the size of each partition balanced according to some criterion. Usually this criterion is described as a percentage of allowed imbalance. Secondly, it needs to minimize the amount of edges that strides across partitions, or in other words, minimize the *edge cut* between partitions. The endpoints of such edges is called a *separator* element. Formally, a separator S in a graph G is a set of elements such that $G \setminus S$ is a disconnected graph.

There are no fixed criterion of the imbalance value that leads to optimal communication reduction. It is possible to assign the entire graph to one part, leaving the other parts of the partitions empty, which effectively eliminates all communication, but this is clearly not optimal as the entire computational workload is assigned to a single process, leaving all other processes with no work. Usually an imbalance value of around 3% is sufficient for most matrices.

Chapter 4

Communication Strategies

This thesis evaluates and compares several communication strategies for Sparse Matrix-Vector Multiplication (SpMV) in a parallel, distributed memory setting. During each iteration of SpMV, every process computes a partial result of the output vector y .

In subsequent iterations, these computed values may be required by other processes to proceed with their own calculations. To ensure correctness, it is therefore necessary to communicate between processes so that each has access to the values it depends on. This section outlines a progression of increasingly efficient strategies for managing this communication.




Figure 4.1: examplegraph

4.1 Exchange entire vector

The most straightforward approach is to have each rank send all of its computed values of y to every other rank. This ensures that all processes possess a complete and updated copy of the output vector before the next iteration. This strategy can be implemented using MPI's collective communication operation `MPI_Allgatherv`, which accommodates variable message sizes from each rank. Figure 4.2 illustrates the state of the y vector before and after communication using this strategy.




Figure 4.2: 1acomm 2

Algorithm 3: 1a - Exchange entire vector

```

for each iteration do
    spmv(g,x,y)
    MPI_Allgatherv(local_y, sendcount, MPI_DOUBLE, y,
        recvcunts, displs, MPI_DOUBLE, MPI_COMM_WORLD)
    swap pointers of  $x$  and  $y$ 

```

4.2 Exchange only separators

An improvement to the previous strategy can be achieved by recognizing that only separator values—those required by multiple processes—must be communicated. Non-separator values are used exclusively by the process that computed them and therefore do not need to be communicated.

To facilitate this strategy, separator values are reordered such that they appear at the beginning of each process's local segment of y . Once this

structure is established, communication is performed using `MPI_Allgatherv`, transmitting only the subset of y that contains separator values. The number of separators on each process must be known beforehand, which can be computed by counting the number of elements that have neighbours belonging to a different partition.




Figure 4.3: 1bcomm

4.2.1 Reordering

After partitioning the matrix into different parts, we obtain a partition vector p , where the $p[i]$ stores the index of the partition the i^{th} entry in A_p . It is necessary to reorder the entries in A_p in accordance with the partition vector, such that all entries belonging to the same partition are stored in sequence. The algorithm below gives an outline of how this can be achieved. Here n_p is the number of partitions, n_r is the size of A_p , and n_c is the size of A_j and A_x .

Algorithm 4: Reordering of Separators

Input : $n_p, n_r, n_c, p, A_p, A_j, A_x$
Output : Reordered A_p, A_j, A_x
 $\text{newId} \leftarrow [0] \cdot n_r$
 $\text{oldId} \leftarrow [0] \cdot n_r$
 $\text{id} \leftarrow 0$
 $p_0 \leftarrow 0$

for $r \in \{0, \dots, n_p - 1\}$ **do**
 for $i \in \{0, \dots, n_r - 1\}$ **do**
 if $p[i] = r$ **then**
 $\text{oldId}[\text{id}] \leftarrow i$
 $\text{newId}[i] \leftarrow \text{id}$
 $\text{id} \leftarrow \text{id} + 1$
 $p[r+1] \leftarrow \text{id}$

$\text{newV} \leftarrow [0] \cdot (n_r + 1)$
 $\text{newE} \leftarrow [0] \cdot n_c$
 $\text{newA} \leftarrow [0] \cdot n_c$

for $i \in \{0, \dots, n_r - 1\}$ **do**
 $\text{degree} \leftarrow A_p[\text{oldId}[i] + 1] - A_p[\text{oldId}[i]]$
 $\text{newV}[i+1] \leftarrow \text{newV}[i] + \text{degree}$

for $i \in \{0, \dots, n_r - 1\}$ **do**
 $\text{degree} \leftarrow A_p[\text{oldId}[i] + 1] - A_p[\text{oldId}[i]]$
 for $j \in \{0, \dots, \text{degree} - 1\}$ **do**
 $\text{newE}[\text{newV}[i] + j] \leftarrow A_j[A_p[\text{oldId}[i]] + j]$
 $\text{newA}[\text{newV}[i] + j] \leftarrow A_x[A_p[\text{oldId}[i]] + j]$
 for $j \in \{\text{newV}[i], \dots, \text{newV}[i+1] - 1\}$ **do**
 $\text{newE}[j] \leftarrow \text{newId}[\text{newE}[j]]$

Overwrite $A_p \leftarrow \text{newV}, A_j \leftarrow \text{newE}, A_x \leftarrow \text{newA}$

4.3 Exchange only required separators

Further reduction to the communication volume can be achieved by observing that not all separator values are required by every process. As the number of partitions increases, the set of dependencies between partitions tends towards sparsity. As a consequence of this, certain sets

of separators may only need to be communicated to a given subset of processes. Using this strategy, each process only communicates its set of separator values to the processes that require them.




Figure 4.4: 1ccomm

4.4 Exchange only required separator values

The final strategy aims to minimize communication overhead by transmitting only the exact subset of separator values that are both computed by and required for inter-process computation. If a specific separator value computed by one process is needed by exactly one other process, then only that single recipient receives the value.

This approach eliminates all unnecessary data transfers but introduces additional complexity in managing communication schedules. Dependencies must be mapped at a fine-grained level, and communication patterns must be explicitly tailored to the structure of the matrix and its partitioning.

4.5 2d - Fully scalable communication




Figure 4.5: 1dcomm

Chapter 5

Results




Figure 5.1: Stacked bar chart of SpMV and Halo Time.




Figure 5.2: Aggregated Results of SpMV on defq chip

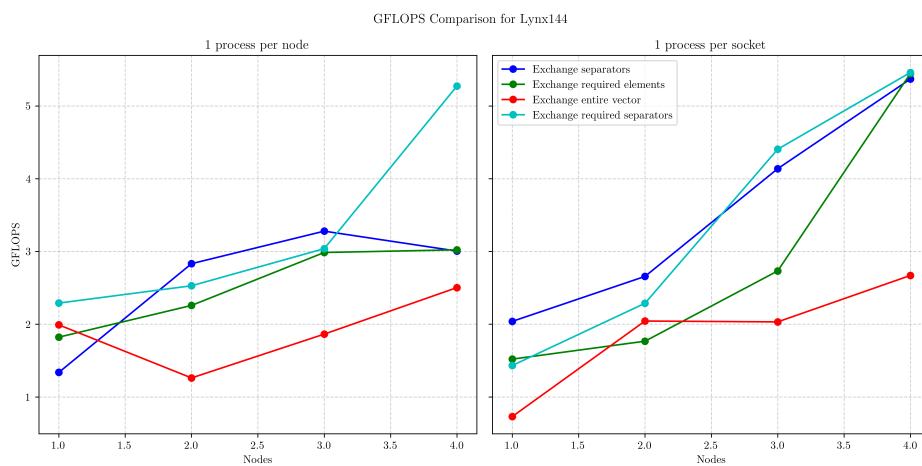


Figure 5.3: SpMV performance with 1 process vs 2 processes per node on Lynx144




Figure 5.4: SpMV performance with 1 process vs 2 processes per node on delaynay_n24




Figure 5.5: SpMV performance with 1 process vs 2 processes per node on rgg_n_2_24_s0

Table 5.1: Hardware used in our experiments. STREAM Triad was run with the `-march=native` and `-O3` compilation flags.

	Xeon-A	Xeon-B	Naples	Rome	Milan	TX2	Hi1620
CPU	Intel Xeon Gold 6130	Intel Xeon Platinum 8168	AMD Epyc 7601	AMD Epyc 7302P	AMD Epyc 7763	Cavium TX2 CN9980	HiSilicon Kunpeng 920-6426
Instr. set	x86-64	x86-64	x86-64	x86-64	x86-64	ARMv8.1	ARMv8.2
Microarch.	Skylake	Skylake	Zen	Zen 2	Zen 3	Vulcan	TaiShan v110
Sockets	2	2	2	1	2	2	2
Cores	2×16	2×24	2×32	1×16	2×64	2×32	2×64
Freq. [GHz]	1.9–3.6	2.5–3.7	2.7–3.2	1.5–3.3	2.5–3.5	2.0–2.5	2.6
L1I/core [KiB]	32	32	64	32	32	32	64
L1D/core [KiB]	32	32	32	32	32	32	64
L2/core [KiB]	1024	1024	512	512	512	256	512
L3/socket [MiB]	22	33	64	16	256	32	64
Mem. channels	2×6	2×6	2×8	1×8	2×8	2×8	2×8
Bandwidth [GB/s]	256	256	342	204.8	409.6	342	342
Triad [GB/s]	147.1	137.4	169.7	90.9	256.5	236.4	260.4




Figure 5.6: SpMV performance with 1 process vs 2 processes per node on road_usa

Chapter 6

Previous Work

Bibliography

- [1] Gautam Gupta, Sivasankaran Rajamanickam, and Erik G. Boman. GAMGI: Communication-reducing algebraic multigrid for gpus. In *Proceedings of the 28th ACM SIGPLAN Annual Symposium on Principles and Practice of Parallel Programming (PPoPP)*, pages 61–75. ACM, 2024.
- [2] Peter Norvig. Latency numbers every programmer should know. <https://norvig.com/21-days.html#Latency>, 2021. Accessed: 2025-04-29.
- [3] Cong Zheng, Shuo Gu, Tong-Xiang Gu, Bing Yang, and Xing-Ping Liu. Biell: A bisection ellpack-based storage format for optimizing spmv on gpus. *Journal of Parallel and Distributed Computing*, 74(7):2639–2647, 2014. Special Issue on Perspectives on Parallel and Distributed Processing.